

# Hypothesis testing

## **Session 2**

MATH 80667A: Experimental Design and Statistical Methods  
for Quantitative Research in Management  
HEC Montréal

# Outline

**Variability**



**Hypothesis tests**

**Power**

**R examples**


# Signal versus noise

# The signal and the noise



Can you spot the differences?

# Sampling variability



# Hypothesis tests

# The general recipe of hypothesis testing

1. Define variables
2. Write down hypotheses (  $\mathcal{H}_0$  /  $\mathcal{H}_a$  )
3. Choose and compute a test statistic
4. Compare the value to the null distribution (benchmark)
5. Compute the  $p$ -value
6. Conclude (reject/fail to reject)
7. Report findings

# Hypothesis tests versus trials

 Scene from "12 Angry Men" by Sidney Lumet



## Trial

- Binary decision: guilty/not guilty
- Summarize evidences (proof)
- Assess evidence in light of **presumption of innocence**
- Verdict: either guilty or not guilty
- Potential for judicial mistakes



# Impact of encouragement on teaching

From Davison (2008), Example 9.2

In an investigation on the teaching of arithmetic, 45 pupils were divided at random into five groups of nine. Groups A and B were taught in separate classes by the usual method. Groups C, D, and E were taught together for a number of days. On each day C were praised publicly for their work, D were publicly reprimanded and E were ignored. At the end of the period all pupils took a standard test.

# Exercise

In pairs, identify

- the experimental and observational units
- the treatment levels
- the response variable
- the null and alternative hypothesis

03 : 00

## Load data

## Summary statistics

## Plot

```
# Load libraries
library(tidyverse)
# Load and reformat data
arithmetic <-
  read_csv("data/arithmetic.csv",
           col_types = c("fn"))
# categorical == factor <fct>
glimpse(arithmetic)
```

```
## Rows: 45
## Columns: 2
## $ group <fct> A, A, A, A, A,...
## $ y      <dbl> 17, 14, 24, 20...
```

# Pick a test, compute its value


One-way analysis of variance uses an  $F$  statistic.

```
#one way analysis of variance  
aov(data = arithmetic,  
     formula = y ~ group)  
# can replace 'aov' by 'lm'
```

- In **R**, the function `anova` prints the analysis of variance table.
- The value of the statistic is 15.268.

How 'extreme' is this number?

# Assessing evidence




## Benchmarking

- The same number can have different meanings
  - units matter!
- Meaningful comparisons require some reference

# Possible, but not plausible

The null distribution tells us what are the plausible values for the statistic and there relative frequency

- what can we expect to see **by chance** if there is **no difference** between groups.




# P-value

Null distributions are different, which makes comparisons uneasy.

- The *P*-values gives the probability of observing an outcome as extreme **if the null hypothesis was true**.

```
pf(stat,  
   df1 = 4,  
   df2 = 40,  
   lower.tail = FALSE)
```



# Power



# I cried power!

The null alternative corresponds to a single value:  $\mu_a = \dots = \mu_e$ .

whereas there are infinitely many alternatives...

Power is the ability to detect when the null is false, for a given alternative.

